

Claudia Platform Revamp

From a horizontal SEO tool suite to one high-value organic growth operator for SaaS companies and digital brands.

Core decision

Claudia will not modify customer codebases directly in the near term. For technical work she produces a copy-ready prompt for Claude Code, Codex, or Cursor. For non-code work she produces exact manual steps. She then asks the owner to confirm deployment and rechecks the result.

"One employee" describes the completed growth loop, not the number of tools in the sidebar.

Prepared for [seogeoaeo.ai](#)

July 2026 | Phase-wise execution plan

01 Executive direction

The product gets stronger by hiding tools and completing the loop.

Revised proposition	Claudia is an organic growth operator for SaaS companies and digital brands. She learns the business, chooses the highest-value SEO, AEO, or GEO action, prepares what needs to ship, and follows up on the result.
----------------------------	---

What changes

- The customer manages Claudia, not separate topics, articles, audits, reports, tools, and scores.
- SEO, AEO, and GEO remain internal methods. They do not become separate product modules.
- Metrics move from the center of the product into evidence attached to a decision or outcome.
- New articles stop being the default. Existing-page refreshes, CTR fixes, answer gaps, and internal links can outrank new content.
- Technical findings become implementation handoffs: coding-agent prompts or manual step-by-step guides.
- Paid tiers unlock the same Claudia. Price changes the amount and pace of work, not intelligence or autonomy mode.

Product doctrine

One visible priority Claudia explains what she is doing now, why it matters, what she will deliver, and what happens after it ships.	Evidence on demand Scores, audits, prompt tracking, and technical details appear only when they support a recommendation or result.
Work before metrics The home screen prioritizes current work and decisions. Counts and composite scores are not the headline value.	Verification is mandatory Every meaningful recommendation must include a baseline, implementation check, and later outcome check.

02 Phase 0 - Great subtraction

Status: implemented in the current product pass.

PHASE 0 Simplify the customer operating surface	Done now Make Claudia feel like the product before adding new agent intelligence.
--	---

Changes completed

- Primary navigation reduced to Claudia, Inbox, Work, and Settings/Account.
- Topics, Articles, Visibility, Reports, Activity, and Tools removed from primary navigation. Their engines and detailed routes remain available from relevant work artifacts.
- Workshop removed from normal navigation, command search, and Settings.
- Dashboard tabs and KPI grid removed. Claudia's current work panel is now the home experience.
- A unified Work destination added over the durable event history. Detailed work types remain evidence, not separate products.
- Public visibility-score snapshot, connector showcase, and redundant pricing design removed from the main marketing path.
- Pricing rebuilt around monthly work capacity and article-draft equivalents.
- The unsupported annual discount removed because checkout does not have an annual price path.
- Prepared fixes now include a coding-agent prompt and a manual implementation guide.

What deliberately remains

Capability	Why it stays	Where the customer sees it
Audit and analyzers	They supply evidence and fix payloads.	Inside a Claudia work item or prepared fix.
Topic and article data	They power planning, drafting, and publishing.	As planned work and produced artifacts.
Visibility and answer tracking	They verify discovery and citation changes.	As evidence and follow-up results.
Reports and activity records	They create trust and an audit trail.	Work history and weekly outcome memo.

Acceptance test	A new customer should understand what Claudia is doing and what she needs from them without learning the names of the underlying analyzers.
------------------------	---

03 The target customer experience

One operator surface with a durable work and evidence system underneath.

Primary destinations

Destination	Purpose	What belongs here
Claudia	Manage the operator.	Mission, current priority, why now, deliverable, next task, weekly outcome summary, and one command surface.
Inbox	Resolve genuine blockers.	Brand judgment, approval, missing connection, permission escalation, or confirmation that a change is live.
Work	See the complete work record.	Planned, active, waiting, and completed work across content, fixes, audits, research, and verification.
Settings	Provide context and guardrails.	Brand, buyer profiles, competitors, work rules, connections, billing, and notifications.

Unified work item

- Objective: the business or organic-growth outcome this work supports.
- Type: refresh, new page, article, metadata, internal links, schema, technical instruction, answer gap, research, or verification.
- Why now: the evidence and reasoning that made this the current priority.
- Expected outcome: what should change and the confidence level.
- Deliverable: draft, file, prompt, manual checklist, comparison, or report.
- Owner action: approve, implement, publish, connect, decide, or no action.
- Verification: immediate live check plus delayed GSC, analytics, or AI-answer follow-up.
- Learning: keep, improve, retry, reverse, or change strategy.

Important boundary	Detailed entity pages can continue to exist for editing and evidence. They are destinations reached from Work, not permanent top-level jobs the customer must operate.
---------------------------	--

04 Implementation handoff model

Deliver value today without direct access to customer codebases.

Technical or code-related work

- Claudia explains the issue, severity, business reason, affected page, and generated artifact.
- She generates one repository-aware prompt designed for Claude Code, Codex, or Cursor.
- The prompt instructs the coding agent to inspect the repository first, follow existing conventions, make the smallest safe change, avoid unrelated refactors, and run the narrowest relevant validation.
- The coding agent must report files changed, reasoning, checks run, and the exact deployed verification URL or command.
- The owner returns to Claudia, marks the work installed, and Claudia schedules a recheck.

Non-code or CMS work

- Claudia identifies the correct admin area: CMS editor, SEO plugin, hosting file manager, site builder, or analytics setting.
- She provides numbered steps, the exact content to paste, a backup instruction, and a live verification step.
- Guides are specific by fix type: schema, metadata, robots.txt, llms.txt, answer blocks, or performance escalation.
- The final step always returns the owner to Claudia for confirmation and later evidence follow-up.

Required handoff states

State	Meaning	Customer action
Prepared	Claudia has produced the prompt, instructions, and artifact.	Copy the appropriate handoff.
Implementation pending	The change is outside Claudia's environment.	Apply it in code, CMS, hosting, or analytics.
Installed	The owner confirms the change is live.	Provide a URL or completion confirmation.
Verified	Claudia confirms the live implementation.	No action unless evidence differs from expected.
Outcome pending	Enough time has not passed for performance evidence.	Wait for the scheduled follow-up.
Learned	The result is classified and the plan is updated.	Review only if Claudia recommends a strategic change.

05 Capacity-based pricing

Do not price Copilot and Autopilot differently when their underlying work cost is the same.

Pricing rule	Every paid tier includes the same Claudia, skills, evidence, prompt handoffs, publishing workflows, and weekly reporting. Plans only change work credits, daily pace, monitoring volume, and supported brand scale.
--------------	---

Plan	Price	Credits	Article equivalent	Pace	Monitoring
Indie	\$29/mo	2,000	Up to 20	1 draft/day	Monthly; 5 questions; 1 competitor
Startup	\$69/mo	5,000	Up to 50	3 drafts/day	Monthly; 10 questions; 3 competitors
Scale	\$199/mo	22,000	Up to 220	10 drafts/day	Weekly; 25 questions; 5 competitors
Enterprise	\$499/mo	130,000	Up to 1,200	40 drafts/day	Weekly; 100 questions; 20 competitors

How to explain credits

- Call them work credits in customer-facing copy. They represent how much heavy work Claudia can complete.
- Use article-draft equivalents only as a comparison unit. State that the maximum assumes all credits are spent on writing.
- Actual output is a mix chosen by priority: articles, research, audits, answer checks, competitor work, and prepared fixes.
- Do not promise a fixed number of articles if audits and research draw from the same capacity.
- Keep top-ups for unusual bursts. Purchased top-up credits do not expire.
- Do not display an annual discount until an annual Stripe price and checkout path actually exist.

Current action costs

Action	Credits	Customer meaning
Article generation	100	One research-backed long-form draft.
Research run	20	A fresh ranked opportunity set.
Visibility audit	50	A new site, search, and AI-readiness evidence pass.
AI answer check	10	One answer-share tracking run.
Competitor discovery	15	Refresh the candidate competitor set.
Competitor benchmark	40	A deeper evidence comparison.

06 Phase 1 - One work operating system

Unify the visible product without forcing an immediate database rewrite.

PHASE 1 Create the Work read model	1-3 weeks Topics, content, findings, tasks, and reports feel like one queue.
---	--

Build

- Define a shared WorkItem view model over the existing task, event, topic, article, finding, approval, and report records.
- Add filters for planned, active, waiting, and completed. Type filters stay secondary.
- Attach evidence, expected outcome, artifact, owner action, verification state, and learning to each row.
- Make detailed article, visibility, and fix pages children of Work through contextual links and breadcrumbs.
- Turn the weekly report into a timeline event and Claudia home summary; keep the archive route as a detail view.
- Keep a durable action ledger for every publish, prompt handoff, manual confirmation, and recheck.

Do not build yet

- No new social, backlink, local SEO, ecommerce, or agency CRM features.
- No rewrite of every underlying table before the customer-facing Work model proves useful.
- No new composite score, dashboard, or standalone analyzer page.

Acceptance criteria

Comprehension A first-time user can describe what Claudia is doing now, why, and what happens next within 30 seconds.	Navigation No routine task requires choosing between Topics, Articles, Visibility, Reports, Activity, and Tools.
Traceability Every visible work item links to its artifact and supporting evidence.	Inbox quality Inbox contains only decisions or implementation confirmations, not ordinary status updates.

07 Phase 2 - Opportunity planner

Replace content-volume scheduling with mixed-task prioritization.

PHASE 2

Choose the highest-value work

3-6 weeks

New articles compete fairly with refreshes, CTR fixes, answer gaps, links, and technical handoffs.

Opportunity types

Existing-page growth Striking-distance query, CTR improvement, decay recovery, content refresh, consolidation, or conversion-path improvement.	Answer visibility Buyer question not citing the brand, weak quotable passage, missing direct answer, or source-authority gap.
Site structure Internal linking, metadata, schema, crawlability, indexation, or page-template issue.	Net-new demand Missing use-case, comparison, integration, glossary, product, or supporting article with defensible demand.

Ranking model

Recommended conceptual score:

Priority	Expected business impact x confidence x strategic relevance, divided by effort, risk, and owner dependency.
-----------------	---

- Impact uses buyer intent, current impressions, ranking distance, traffic value, conversion relevance, and answer visibility.
- Confidence uses evidence freshness, agreement across sources, data volume, and similarity to past wins.
- Effort uses credits, generation time, implementation complexity, and review requirements.
- Risk uses brand sensitivity, technical blast radius, uncertainty, and reversibility.
- Owner dependency penalizes work that cannot move without missing access, context, or approval.

Planner behavior

- Select one active priority and at most two visible next items.
- Explain why the winner outranked the alternatives.
- Prefer improving an existing relevant page when expected impact exceeds a new page.
- Replan after new evidence, completion, rejection, failed implementation, or outcome learning.
- Use deterministic policy for permissions and budget. Use the model for interpretation, planning, and explanation.

08 Phase 3 - Existing content and implementation quality

Make the work more valuable before increasing volume.

PHASE 3 Refresh and handoff excellence	4-8 weeks Claudia improves the current site as confidently as she writes new content.
---	---

Existing-content inventory

- Build a canonical inventory of indexable pages, page type, intent, primary queries, impressions, clicks, position, internal links, and last meaningful update.
- Detect decay, cannibalization, thin pages, duplicated intent, missing buyer-stage coverage, orphan pages, and weak internal-link paths.
- Create refresh briefs that preserve useful content, specify additions and removals, and state the expected search or answer outcome.
- Track whether a refresh, consolidation, or new page created the result. Do not merge all content output into one article count.

Handoff quality standard

Deliverable	Must include	Verification
Coding-agent prompt	Repo inspection, exact issue, artifact, constraints, validation commands, reporting format.	Changed files, checks run, deployed URL, source or test evidence.
Manual guide	Admin location, backup, numbered steps, exact value, publish instruction.	Live URL, source check, tool recheck, owner confirmation.
Content refresh	Query thesis, outline delta, facts to preserve, links, title/meta, conversion path.	Live copy, indexability, internal links, later GSC checkpoint.
New content	Buyer intent, gap evidence, unique thesis, brand voice, citations, distribution path.	Publish state, indexability, impressions, clicks, assisted conversion.

Quality gate	Claudia should produce fewer work items if that is what it takes to make every handoff implementable, brand-correct, and verifiable.
---------------------	--

09 Phase 4 - Verification and learning

Turn completed activity into proof and better future decisions.

PHASE 4 Verify every action	4-8 weeks Each work item has an immediate live check and a delayed outcome check.
---	---

Verification sequence

Checkpoint	Question	Evidence
Baseline	What is true before the change?	Current HTML, audit result, GSC query/page data, AI answer, or analytics event.
Installed	Did the requested implementation happen?	Owner confirmation, deployed URL, source response, schema/robots/meta check.
Live verified	Is the change technically correct and accessible?	Fresh crawl, source inspection, response status, structured-data validation.
Early signal	Did discovery or engagement begin to move?	Indexation, impressions, AI mention, click-through, or crawl result.
Outcome	Did the target organic result improve?	Day 7/28/90 page and query evidence, high-intent visits, or selected conversion.
Learning	What should Claudia do next?	Keep, expand, refresh, reverse, retry, or choose another strategy.

Customer-visible outcomes

- A buyer-intent query entered the top 10 or improved materially.
- A target page gained qualified clicks or recovered lost demand.
- A tracked buyer question began citing or accurately describing the brand.
- Organic visits reached a selected high-intent page or conversion event.
- A technical issue was implemented, verified live, and stayed resolved on recheck.

Weekly outcome memo

- What Claudia changed or prepared.
- Why the work mattered.
- What evidence moved, stayed flat, or is still too early to judge.
- What Claudia will do next.
- At most one owner ask.

10 Product metrics and anti-bloat guardrails

Measure whether Claudia creates useful outcomes and reduces customer effort.

North star

Verified high-intent organic wins	The number of completed actions that produce a validated improvement in buyer-intent search visibility, qualified organic traffic, AI-answer presence, or a selected conversion signal.
--	---

Supporting metrics

Time to first useful work Time from brand creation to the first evidence-backed recommendation and usable deliverable.	Implementation rate Share of prepared prompts or manual fixes confirmed live within 14 days.
Verification coverage Share of completed work with a stored baseline, live check, and scheduled outcome checkpoint.	Owner effort Decisions requested per verified win and median time spent in the product each week.
Planner quality Share of top-ranked work accepted, implemented, and later classified as useful.	Retention ritual Weekly outcome memos opened and work resumed after a meaningful result or blocker.

Do not use as headline success

- Composite visibility score without supporting business evidence.
- Number of topics queued, articles generated, prompts tracked, or audits run.
- Credit consumption by itself. Spending is capacity usage, not customer value.
- Number of integrations or standalone tools available.

Feature gate

Before adding a feature, ask	Required answer
Does Claudia need this to complete the current organic-growth loop?	Yes, and the missing step is common in the selected market.
Can this be an internal tool instead of a customer-facing page?	Default to internal unless the customer has a recurring decision to make.
Which existing surface owns it?	Claudia, Inbox, Work, or Settings. No new primary destination without exceptional evidence.
How will the action be verified?	A baseline, live check, outcome checkpoint, and learning rule are defined before launch.
What will be removed or simplified?	Every meaningful addition should reduce another manual step or obsolete an older surface.

11 90-day sequence and future boundary

Complete the operator loop before expanding the feature catalog.

Window	Product focus	Definition of done
Days 0-15	Stabilize the subtraction pass and collect usability evidence.	Three-route navigation holds; no KPI dashboard regression; pricing copy matches checkout and capacity rules.
Days 15-35	Build the unified Work read model and implementation states.	All active content, fix, audit, and approval work appears in one status model with artifact and evidence links.
Days 35-65	Ship the mixed-task opportunity planner and existing-page inventory.	Claudia ranks refreshes, fixes, answer gaps, links, and new content in one queue with explainable reasoning.
Days 65-90	Complete verification and weekly outcome storytelling.	Installed work gets a live recheck, delayed checkpoint, result classification, and concise weekly memo.

Explicitly deferred: direct codebase and site execution

Future phase only	WordPress write actions, GitHub pull requests, rollback, and automatic live-site changes remain future work. Revisit only after prompt/manual handoffs show strong implementation demand, verification is reliable, and customers are asking to remove that final step.
--------------------------	---

Do not add before the core loop is proven

- Backlink marketplace or outreach automation.
- General social media management.
- Local SEO suite.
- Ecommerce-specific optimization suite.
- Agency prospecting, proposals, or white-label dashboards.
- More standalone analyzers or composite scores.
- Unimplemented network integrations shown as available product surface.

Final product test

The customer should be able to say:	"Claudia knows my brand, tells me the best organic-growth action, gives me exactly what I need to ship it, and comes back with the result."
--	---